

Text Analysis

The case of the unnamed .txt files...

James Murray

May '21



The 42 text files have been successfully grouped into 6 different topics. These topics are somewhat related as the analysis below will demonstrate. Folder names have been suggested for each of the topic folders in which the files have been allocated to. There is significant opportunity to extend this sort of analysis into other business areas while understanding our own organisational culture more thoroughly. This can assist in reducing risk and identifying friction points ahead of time.

Contents

Business Understanding	3
Objectives and Assessment of the Situation	3
Data Mining Goals and Project Plan	4
Describe the Data	4
Exploratory Analysis	5
Modeling	12
Evaluation	13
Next Steps	14
Future Opportunities	14
Annex A - Suggested File Groupings	15
Annex B - Suggested file names	18
References	20

List of Tables

1	Word and Sentence counts for the files	4
2	Most frequently occurring words in file 39	5
3	Most Frequently occurring terms in file 29	5
4	Topic 1 contents and confidence levels	15
5	Topic 2 contents and confidence levels	15
6	Topic 3 contents and confidence levels	15
7	Topic 4 content and confidence levels	16
8	Topic 5 contents and confidence levels	16
9	Topic 6 contents and confidence levels	17

List of Figures

1	WordCloud highlighting most used terms across all files	5
2	Top Terms according to Zipf's Law	6
3	Term Frequency for file 08	7
4	Term Frequency for file file 39	8
5	Bi-gram Frequency for file 01	9
6	Bi-gram Frequency for file 38	10
7	Bi-gram network relationship across all files	11
8	Top 10 Terms by topic following LDA modelling	12
9	Classification of Documents by Topic	13

Business Understanding

A series of *random* files have been found in a folder on a personal computer, these files are unhelpfully named, 01.txt - 42.txt files. They are of varying lengths, the relative importance of each file isn't known. To go through the files is both time consuming and arduous - it could be a trip down memory lane to rediscover some old work but time poor folks don't want to *waste* time doing this. Ideally these files would all be grouped into similar topics where appropriate and filed under these topics, maybe even renamed with key terms to make future referencing simple.

Objectives and Assessment of the Situation

The situation isn't dire, there isn't any time pressure on and given that these files have not been directly searched for implies this will be a house keeping activity. The final outcome sought will be a series

of folders with topic titles and in these folders a set of named documents that will be more readily searchable for future referencing.

Data Mining Goals and Project Plan

The Data has been given in a single folder named **Files**. There is no requirement to search for more data beyond this file. The broad plan is to conduct a text analysis on the files to determine some broad topics that these files can be grouped into. The data is a folder with 42 separate text files (.txt) in it.

Describe the Data

The summary table below offers an initial indication of the numbers of words and sentences in each of the files.

Table 1: Word and Sentence counts for the files¹

Item	Minimum	Mean	Maximum
Words	508	2235	5785
Sentences	32	100	263

¹ This is inclusive of numbers and special characters for an initial assessment. Sentences have been assessed in accordance with unicode sentence boundaries. (UAX #29, n.d.)

Below is a *wordcloud* that highlights the most used terms across all of the files.



In addition to the *wordcloud* a list most frequently occurring words for each file was developed, an example of the tables generated is below:

word	n
correlated	35
correlation	31
distribution	25
duration	22
tasks	20
uncorrelated	20

word	n
management	60

word	n
project	44
work	37
bureaucratic	27
post	19

A table was generated for each file which offers insight into the contents of the files.

Another way to start to consider the importance of each word in the files is to use apply Zipf's law to the corpus of documents. Below is a broad list of the terms that are considered important to each of the files and the relative importance according to Zipf's Law.²




Figure 2: Top Terms according to Zipf's Law

² Zipf's is an empirical law stating that word frequency in documents is inversely proportional to work rank in descending order of occurrence. ((PDF) Zipf's Law and Heaps' Law Can Predict the Size of Potential Words, n.d.)

This was done in addition to producing a graph for each of files that highlighted the most frequent terms in each file. Two examples are below.




Figure 3: Term Frequency for file 08




Figure 4: Term Frequency for file file 39

In order to get a better understanding of the contents of the files the files were broken into N-grams or groupings of words that are used together in the documents. This offers a way to gain greater insight by understanding which words appear together or which words appear in the same document generally.(4 Relationships between words, n.d.) Example results for two of the files are below.

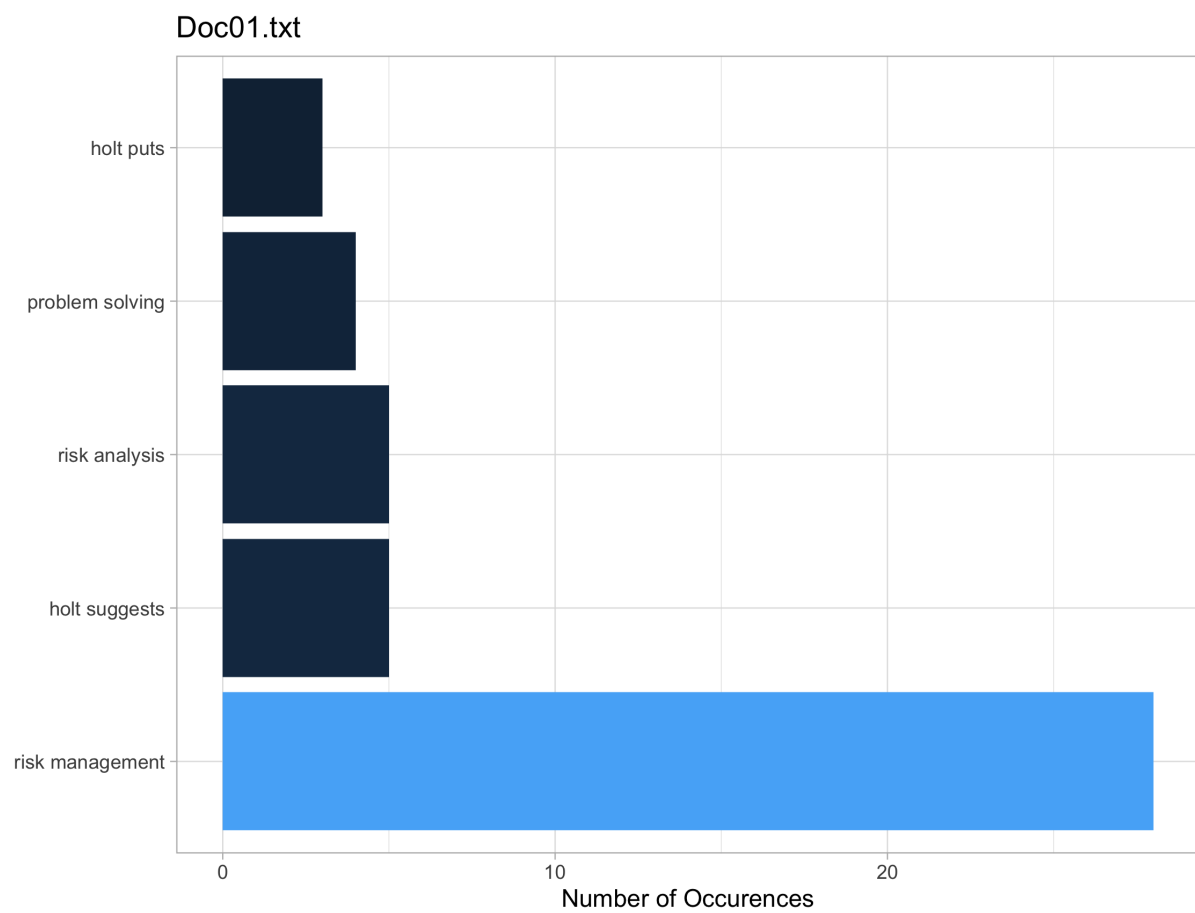


Figure 5: Bi-gram Frequency for file 01




Figure 6: Bi-gram Frequency for file 38

From visual analysis we can see that this is more informative but we still don't have a really good feel for how many different possible topics there are across all of the files or what reasonable groupings for the documents could be. considering the terms closely it is evident that some more cleaning of the text corpus is required. This will be addressed as part of the steps before unsupervised modelling is conducted.

The relationships across all of the documents were mapped visually in an effort to better understand groupings of the different terms across all of the files. This relationship diagram is below.




Figure 7: Bi-gram network relationship across all files

Based on this gross check we can see a few key terms that are starting to group together - this can assist in informing the number of topics as part of future modelling activities.

The file preparation to this point has consisted of the removal of digits and a series of stopwords. Some extra *stop words* were added to the list based on initial analysis. Characters like ‘_t’ were removed from the files.

Modeling

Prior to conducting unsupervised topic modelling using the *Latent Dirichlet Allocation*³ all punctuation was removed in addition to *stemming*⁴ of the Data. This comes with some risk as some of the sentiment can be lost as part of the process - it helps to reduce the feature space as part of larger data sets - it was done here to determine if there are any benefits in narrowing the topics.

The results of the topics modelling are below.

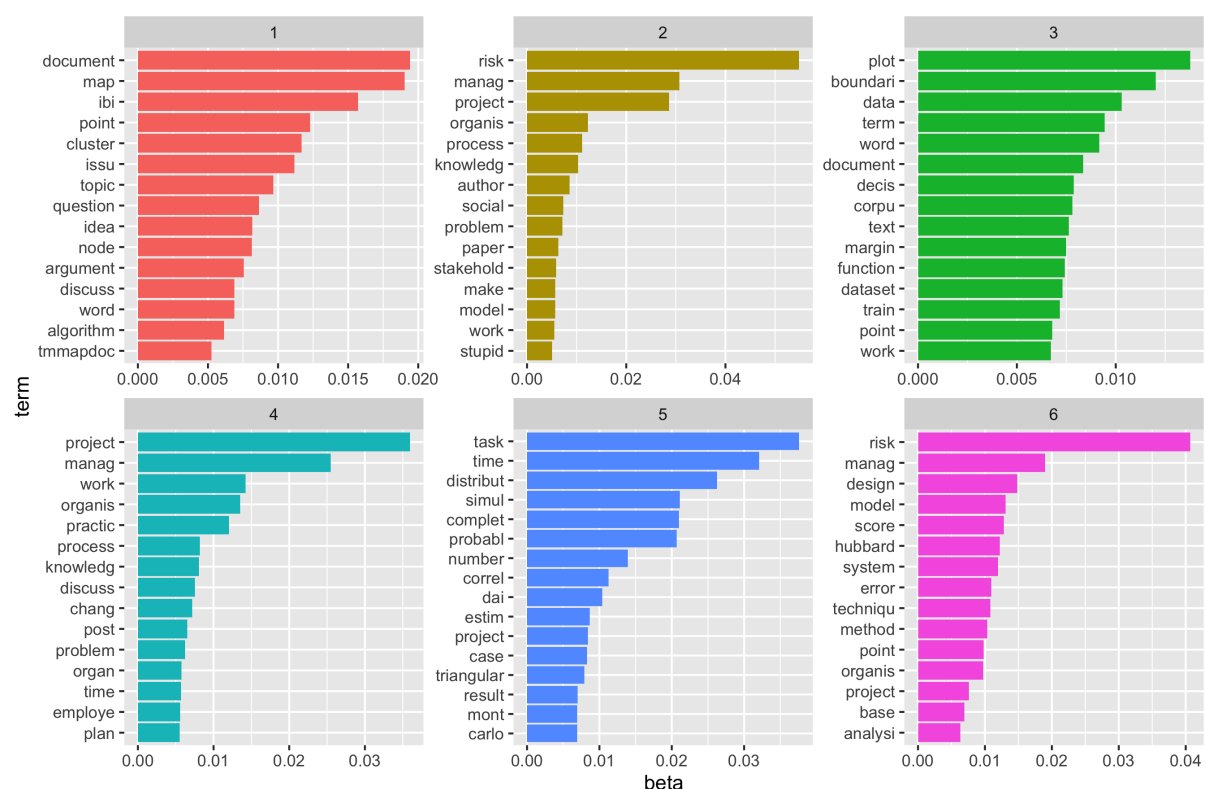


Figure 8: Top 10 Terms by topic following LDA modelling

The choice to make six topics as part of the process was after a series of trials - this number was assessed as offering the best balance between meaningful groups and the size of the groups being useful. There is still some overlap between the terms and the topics but it is likely this will be a product of the content in the files themselves.

In order to assess how the model went at classifying the documents into different groups the facet chart below was created. This offers a visual way to view which documents should be grouped into which topic and how confident we are of the grouping.

³ LDA is one of the most common algorithms for topic modelling. (6 Topic modeling | Text Mining with R, n.d.)

⁴ Stemming removes the different versions of the same word. (Silge, n.d.a)

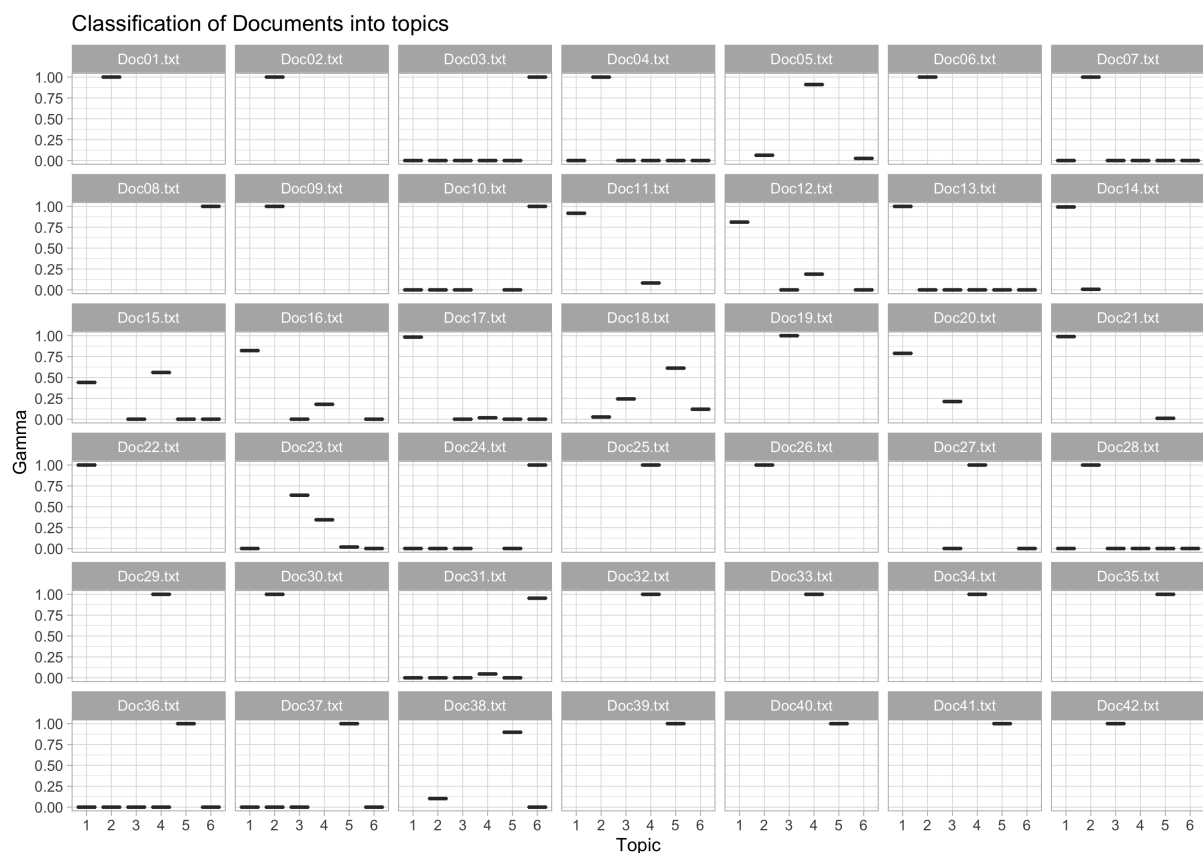


Figure 9: Classification of Documents by Topic

It appears that of the 42 files only a couple of files couldn't be conclusively placed into a topic. These files are 15, 18 and 23. After reviewing these files there is some overlap across topics and the terms used in these files with other topics. File 15 should be grouped into topic 4, file 8 should be grouped into topic 5. File 23 was one of the harder to group. The model had it sitting in topic 5 but this doesn't seem suitable - this file will require further review before settling on a topic as it appears to span multiple topics equally and none of them conclusively. As per Figure 2 above we can see that the key terms in the document are *coal*, *seam*, *mines*, *mechanised* which tends to indicate a unique set of terms according to Zipf's law relative to other documents.

Evaluation

The analysis conducted on the *mystery* files has revealed 6 different topics according to LDA analysis. LDA was conducted with both *stemming* and without it - the results of the model that was preprocessed with *stemming* were marginally better offering a few more files that could be conclusively grouped.

Next Steps

Following this analysis the way forward would be to create 6 different files using terms from the six different topics as a labelling convention, suggested files names and the contents of each file are in Annex A. A list of suggested file names is at Annex B.

Future Opportunities

Given the analysis conducted in this paper there is latent ability within the firm to analyse text to gain greater insight into the language used across a variety of documents. This analysis when combined with other sources of information can be leveraged to better understand both our own language within the firm in addition to the language that other are using when referring to the organisation. Although this was a relatively simple task to understand the contents and group these files for future reference the power of understanding language and how it influences the company and how the company is perceived should not be understated. It will influence and affect everything we do, it can be a barometer for company culture, when combine with other sources can offer a more fulsome view of the organisation in a competitive market place.

Future work opportunities could be to conduct sentiment analysis on our own reviews, analyse our internal communications to understand our company culture, combine sentiment analysis with market movements and trends to determine if there are any lead or lag indicators around pricing opportunities.

Annex A - Suggested File Groupings

Table 4: Topic 1 contents and confidence levels

document	gamma
Doc19.txt	0.9999504
Doc42.txt	0.9999456

Suggested folder name is: IBIS_MAP_CLUSTER

Table 5: Topic 2 contents and confidence levels

document	gamma
Doc35.txt	0.9999317
Doc40.txt	0.9999015
Doc39.txt	0.9998940
Doc41.txt	0.9998747
Doc37.txt	0.9997340
Doc36.txt	0.9997073
Doc38.txt	0.8971244

Suggested folder name is: RISK_MANAGEMENT_STUPID

Table 6: Topic 3 contents and confidence levels

document	gamma
Doc29.txt	0.9999198
Doc25.txt	0.9998651
Doc34.txt	0.9998628
Doc33.txt	0.9998570
Doc32.txt	0.9998552
Doc27.txt	0.9998268

document	gamma
Doc05.txt	0.9103705

Suggested folder name is: PLOT_BOUNDARY_CORPUS

Table 7: Topic 4 content and confidence levels

document	gamma
Doc08.txt	0.9999133
Doc24.txt	0.9997707
Doc10.txt	0.9997646
Doc03.txt	0.9996496
Doc31.txt	0.9535174

Suggested folder name is: PROJECT_MANAGEMENT_CHANGE

Table 8: Topic 5 contents and confidence levels

document	gamma
Doc22.txt	0.9999089
Doc13.txt	0.9996781
Doc14.txt	0.9929273
Doc21.txt	0.9894054
Doc17.txt	0.9830871
Doc11.txt	0.9175054
Doc16.txt	0.8216867
Doc12.txt	0.8116897
Doc20.txt	0.7884268

Suggested folder name is: MONTE_CARLO_PROJECT

Table 9: Topic 6 contents and confidence levels

document	gamma
Doc30.txt	0.9999035
Doc01.txt	0.9998910
Doc09.txt	0.9998767
Doc02.txt	0.9998722
Doc26.txt	0.9998696
Doc06.txt	0.9998590
Doc28.txt	0.9996987
Doc07.txt	0.9996730
Doc04.txt	0.9993338

Suggested folder name is: RISK_MODEL_PROJECT

Annex B - Suggested file names

These names were generated from the terms that are a product of the top five terms from each file as a product of Zipf's Law.

doc_id	file_name_sug
Doc01.txt	holt_pschoanalysis_risk_risks_fortune_Doc01.txt
Doc02.txt	erm_appetite_risk_controls_risks_Doc02.txt
Doc03.txt	risk_risks_historical_pitfalls_module_Doc03.txt
Doc04.txt	elephant_elephants_elephants_dead_pachyderms_Doc04.txt
Doc05.txt	schedule_signs_percent_trouble_ewss_Doc05.txt
Doc06.txt	social_risks_risk_stakeholders_construct_Doc06.txt
Doc07.txt	strategic_risks_authors_deliverable_operational_Doc07.txt
Doc08.txt	hubbard_risk_analysts_models_risks_Doc08.txt
Doc09.txt	risk_risks_authors_interviewee_management_Doc09.txt
Doc10.txt	scoring_matrices_risk_methods_hubbard_Doc10.txt
Doc11.txt	ibis_glyma_maps_students_video_Doc11.txt
Doc12.txt	ibis_nodes_decision_mapping_compendium_Doc12.txt
Doc13.txt	map_rohan_sleep_bed_movie_Doc13.txt
Doc14.txt	mary_jack_andrew_max_rick_Doc14.txt
Doc15.txt	informal_knowledge_capture_ibis_conklin_Doc15.txt
Doc16.txt	ibis_map_dialogue_mapping_complexity_Doc16.txt
Doc17.txt	infrastructural_carr_map_carrs_stage_Doc17.txt
Doc18.txt	machine_data_learning_fit_model_Doc18.txt
Doc19.txt	corpus_text_wordcloud_freqr_dtm_Doc19.txt
Doc20.txt	cluster_clustering_clusters_documents_corpus_tmmap_Doc20.txt
Doc21.txt	topic_lda_topics_tmmap_document_Doc21.txt
Doc22.txt	gephi_graph_similarity_document_tmmap_Doc22.txt
Doc23.txt	mechanised_mines_seam_shifts_coal_Doc23.txt
Doc24.txt	scapegoat_author_active_errors_error_Doc24.txt
Doc25.txt	willmott_alvesson_critical_employees_lifeworld_Doc25.txt

doc_id	file_name_sug
Doc26.txt	stupidity_functional_employees_spicer_alvesson_Doc26.txt
Doc27.txt	effects_planned_change_side_improvisations_reforms_unintended_Doc27.txt
Doc28.txt	success_npd_leadership_culture_product_variables_Doc28.txt
Doc29.txt	bureaucratic_buzzbank_hodgson_creative_employees_Doc29.txt
Doc30.txt	sharing_knowledge_creation_model_isd_jackson_klobas_Doc30.txt
Doc31.txt	design_conway_paths_communication_subsystems_Doc31.txt
Doc32.txt	grabher_projects_military_hosting_institutions_Doc32.txt
Doc33.txt	practices_cerrits_wareham_recipient_practice_Doc33.txt
Doc34.txt	validity_john_claims_habermas_speaker_Doc34.txt
Doc35.txt	completion_days_ml_max_probability_Doc35.txt
Doc36.txt	dartboard_hits_theta_centre_demo_Doc36.txt
Doc37.txt	dart_circle_square__throws_Doc37.txt
Doc38.txt	leq_risk_probability_hrs_task_Doc38.txt
Doc39.txt	correlated_uncorrelated_correlation_coefficient_duration_Doc39.txt
Doc40.txt	completion_hrs_tasks_task_distribution_Doc40.txt
Doc41.txt	ml_max_simulation_distribution_triangular_Doc41.txt
Doc42.txt	boundary_margin_svm_dataset_svmmodel_Doc42.txt

References

- 1 The tidy text format | Text Mining with R n.d. Available at <https://www.tidytextmining.com/tidytext.html> [Last accessed 24 May 2021].
- 4 Relationships between words: N-grams and correlations | Text Mining with R n.d. Available at <https://www.tidytextmining.com/ngrams.html> [Last accessed 29 May 2021].
- 6 Topic modeling | Text Mining with R n.d. Available at <https://www.tidytextmining.com/topicmodeling.html> [Last accessed 29 May 2021].
- Make ggplot2 purrr n.d. Available at <https://www.brodrigues.co/blog/2017-03-29-make-ggplot2-purrr/> [Last accessed 28 May 2021].
- (PDF) Zipf's Law and Heaps' Law Can Predict the Size of Potential Words n.d. *ResearchGate*. DOI: <https://doi.org/10.1143/PTPS.194.202>
- Reading text files with readtext n.d. Available at https://cran.r-project.org/web/packages/readtext/vignettes/readtext_vignette.html [Last accessed 24 May 2021].
- Silge, E H and J** n.d.b *Supervised Machine Learning for Text Analysis in R*. Available at <https://smlltar.com/> [Last accessed 24 May 2021].
- Silge, E H and J** n.d.a *Chapter 4 Stemming | Supervised Machine Learning for Text Analysis in R*. Available at <https://smlltar.com/stemming.html> [Last accessed 28 May 2021].
- UAX #29: Unicode Text Segmentation n.d. Available at https://www.unicode.org/reports/tr29/#Sentence_Boundaries [Last accessed 28 May 2021].